

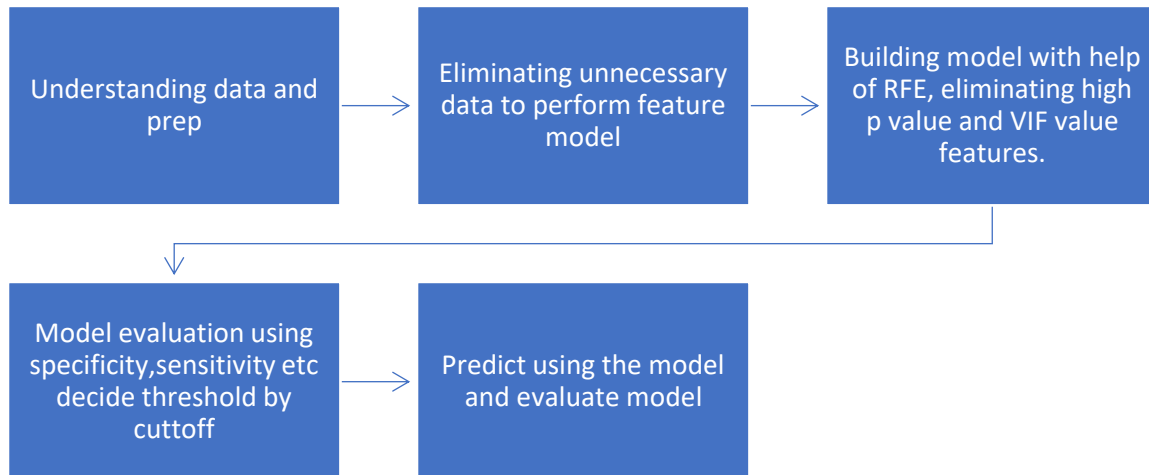
LEAD SCORING CASE STUDY- PPT

TO BUILD A LOGISTIC REGRESSION MODEL TO PREDICT IF A LEAD FOR ONLINE COURSES FOR AN EDUCATION COMPANY AKA X EDUCATION WOULD SUCCESSFULLY BE CONVERTED OR NOT.

BUSINESS OBJECTIVE

- To help X Education to select Hot Leads, i.e. the leads that are most likely to convert into paying customers.
- To build a logistic regression model to assign a lead score value between 0 and 100 to each of the leads which can be used by the company to target potential customers.

PROBLEM SOLVING METHODOLOGY



1. Removing null values

```
In [222]: ▶ leads.isnull().sum()

Out[222]: Prospect ID                0
Lead Number                        0
Lead Origin                        0
Lead Source                        0
Do Not Email                       0
Converted                          0
TotalVisits                        0
Total Time Spent on Website        0
Page Views Per Visit              0
Last Activity                      0
Specialization                     0
What is your current occupation    0
A free copy of Mastering The Interview 0
Last Notable Activity              0
dtype: int64
```

2. Removing columns with high missing values

```
In [208]: ▶ # dropping Lead Profile and how did you hear about X Education coz it has high select values
```

```
In [209]: ▶ leads.drop(['Lead Profile', 'How did you hear about X Education'], axis = 1, inplace = True)
```

3. Removing redundant data

```
In [210]: ▶ # dropping other columns with redundant data
```

```
In [211]: ▶ leads.drop(['Do Not Call', 'Search', 'Magazine', 'Newspaper Article', 'X Education Forums', 'Newspaper',
'Digital Advertisement', 'Through Recommendations', 'Receive More Updates About Our Courses',
'Update me on Supply Chain Content', 'Get updates on DM Content',
'I agree to pay the amount through cheque'], axis = 1, inplace = True)
```

4. Dummy Encoding

```
In [229]: ▶ # Create dummy variables
dummy_1 = pd.get_dummies(leads[['Lead Origin', 'Lead Source', 'Do Not Email', 'Last Activity',
'What is your current occupation', 'A free copy of Mastering The Interview',
'Last Notable Activity']], drop_first=True)

leads = pd.concat([leads, dummy_1], axis=1)
```


```
In [230]: ▶ dummy_spl_1 = pd.get_dummies(leads['Specialization'], prefix = 'Specialization')
#dummy_spl_1 = dummy_spl.drop(['Specialization_Select'], 1)
leads = pd.concat([leads, dummy_spl_1], axis = 1)
```

```
In [231]: ▶ leads = leads.drop(['Lead Origin', 'Lead Source', 'Do Not Email', 'Last Activity',
'Specialization', 'What is your current occupation',
'A free copy of Mastering The Interview', 'Last Notable Activity'], 1)
```

5. Test Train Split

```
In [236]: # Split dataset into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size=0.3, random_state=100)
```

6. Checking for correlation



7. Feature Selection

```
In [243]: # checking for features selected by RFE
list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
Out[243]: [('TotalVisits', True, 1),
('Total Time Spent on Website', True, 1),
('Page Views Per Visit', False, 10),
('Lead Origin_Landing Page Submission', False, 2),
('Lead Origin_Lead Add Form', True, 1),
('Lead Origin_Lead Import', False, 52),
('Lead Source_Direct Traffic', False, 15),
('Lead Source_Facebook', False, 48),
('Lead Source_Google', False, 38),
('Lead Source_Live Chat', False, 30),
('Lead Source_Olark Chat', True, 1),
('Lead Source_Organic Search', False, 37),
('Lead Source_Pay per Click Ads', False, 35),
('Lead Source_Press Release', False, 53),
('Lead Source_Reference', False, 4),
('Lead Source_Referral Sites', False, 39),
('Lead Source_Social Media', False, 61),
('Lead Source_WeLearn', False, 23),
('Lead Source_Welingak Website', True, 1),
('Lead Source_bing', False, 19).]
```

8. Model Building

```
In [245]: # creating a Logistic regression model
```

```
In [248]: # Fitting Logistic Regression model on X_train
X_train_sm = sm.add_constant(X_train)
logm2 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = logm2.fit()
res.summary()
```

Out[248]: Generalized Linear Model Regression Results

Dep. Variable:	Converted	No. Observations:	4481
Model:	GLM	Df Residuals:	4445
Model Family:	Binomial	Df Model:	15
Link Function:	logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-2087.2
Date:	Sun, 08 Mar 2020	Deviance:	4134.4
Time:	13:34:07	Pearson chi2:	4.83e+03
No. Iterations:	22		
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
const	-0.9490	0.803	-1.573	0.116	-2.131	0.233
TotalVisits	10.2343	2.636	3.882	0.000	5.068	15.401
Total Time Spent on Website	4.4045	0.188	23.735	0.000	4.041	4.788
Lead Origin_Lead Add Form	4.2381	0.259	16.363	0.000	3.729	4.744
Lead Source_Olark Chat	1.6324	0.133	12.287	0.000	1.372	1.893
Lead Source_Welingak Website	2.3444	1.038	2.258	0.024	0.310	4.379
Do Not Email_Yes	-1.5177	0.192	-7.892	0.000	-1.895	-1.141
Last Activity_Had a Phone Conversation	1.1713	0.987	1.188	0.235	-0.784	3.106
Last Activity_SMS Sent	1.1787	0.082	14.305	0.000	1.017	1.340
What is your current occupation_Housewife	22.6104	2.45e+04	0.001	0.999	-4.8e+04	4.8e+04
What is your current occupation_Student	-1.1280	0.634	-1.776	0.078	-2.389	0.117
What is your current occupation_Unemployed	-1.2988	0.598	-2.169	0.030	-2.488	-0.125
What is your current occupation_Working Professional	1.2483	0.627	1.992	0.048	0.020	2.476
Last Notable Activity_Had a Phone Conversation	23.0108	2.09e+04	0.001	0.999	-4.09e+04	4.1e+04
Last Notable Activity_Unreachable	2.7870	0.807	3.429	0.001	1.188	4.348
Specialization_Select	-0.3400	0.068	-3.484	0.001	-0.532	-0.148

10. Checking VIF values

```
In [251]: # Making VIF dataframe
vif = pd.DataFrame()
vif['Features'] = X_train.columns
vif['VIF'] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[251]:

	Features	VIF
10	What is your current occupation_Unemployed	4.13
6	Last Activity_Had a Phone Conversation	2.44
12	Last Notable Activity_Had a Phone Conversation	2.43
1	Total Time Spent on Website	2.39
14	Specialization_Select	1.90
2	Lead Origin_Lead Add Form	1.71
3	Lead Source_Olark Chat	1.66
0	TotalVisits	1.63
7	Last Activity_SMS Sent	1.59
11	What is your current occupation_Working Profes...	1.56
4	Lead Source_Welingak Website	1.37

11. Fitting p values and vif values

```
In [261]: logm1 = sm.GLM(y_train,(sm.add_constant(X_train)), family = sm.families.Binomial())
logm1.fit().summary()
```

Out[261]: Generalized Linear Model Regression Results

Dep. Variable:	Converted	No. Observations:	4461
Model:	GLM	Df Residuals:	4448
Model Family:	Binomial	Df Model:	12
Link Function:	logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-2072.6
Date:	Sun, 08 Mar 2020	Deviance:	4145.3
Time:	13:34:08	Pearson chi2:	4.81e+03
No. Iterations:	7		
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
const	0.2371	0.196	1.211	0.226	-0.147	0.621
TotalVisits	10.0121	2.618	3.825	0.000	4.882	15.143
Total Time Spent on Website	4.3957	0.185	23.708	0.000	4.032	4.759
Lead Origin_Lead Add Form	4.2341	0.259	16.364	0.000	3.727	4.741
Lead Source_Olark Chat	1.6321	0.133	12.275	0.000	1.371	1.893
Lead Source_Welingak Website	2.3468	1.038	2.260	0.024	0.312	4.382
Do Not Email_Yes	-1.5182	0.192	-7.891	0.000	-1.895	-1.141

12. Model Evaluation- predicting the data

```
In [264]: # using predict
y_train_pred = res.predict(X_train_sm)
y_train_pred[:10]
```

Out[264]: 8003 0.315577
218 0.151844
4171 0.135876
4037 0.278192
3660 0.959650
207 0.156043
2044 0.143676
6411 0.952580
6498 0.079814
2085 0.981919
dtype: float64

```
In [265]: # Reshaping into an array
y_train_pred = y_train_pred.values.reshape(-1)
y_train_pred[:10]
```

Out[265]: array([0.3155766 , 0.1518439 , 0.13587609, 0.27819235, 0.95965009,
0.1560432 , 0.14367596, 0.95258003, 0.07981364, 0.98191931])

```
In [266]: # Creating new dataframe
y_train_pred_final = pd.DataFrame({'Converted':y_train.values, 'Conversion_Prob':y_train_pred})
y_train_pred_final.head()
```


```
In [267]: # Creating a new column predicted_col with 1 if Paid_Prob > 0.5 else 0

In [268]: y_train_pred_final['predicted_col'] = y_train_pred_final.Conversion_Prob.map(lambda x: 1 if x > 0.5 else 0)
y_train_pred_final.head()
```

Out[268]:

	Converted	Conversion_Prob	predicted_col
0	0	0.315577	0
1	0	0.151844	0
2	1	0.135876	0
3	1	0.278192	0
4	1	0.959650	1

13. Finding optimal cutoff




#here area under the curve of the ROC is 0.88 which shows a good model

14. Optimal probability plot

```
In [283]: ##plotting

prob_cutoff_df.plot.line(x='prob', y=['accuracy', 'sensi', 'speci'])
plt.show()
```



Here when probability thresholds are very low, sensitivity value is very high and specificity is very low. Likewise, for higher probability thresholds, sensitivity values are very low but the specificity values are very high.

High sensitivity indicates that the model will accurately identify nearly all leads that are likely to convert.

15. Checking accuracy and confusion matrix

```
In [285]: # checking accuracy
metrics.accuracy_score(y_train_pred_final.Converted, y_train_pred_final.final_predicted)

Out[285]: 0.7906299036090563

In [286]: # creating conf matrix
confusion_mat_2 = metrics.confusion_matrix(y_train_pred_final.Converted, y_train_pred_final.final_predicted )
confusion_mat_2

Out[286]: array([[1816,  496],
                [ 438, 1711]], dtype=int64)

In [287]: # checking other matrices
TP = confusion_mat_2[1,1] # for true positive
TN = confusion_mat_2[0,0] # for true negatives
FP = confusion_mat_2[0,1] # for false positives
FN = confusion_mat_2[1,0] # for false negatives

In [288]: # Calculating Sensitivity
TP/(TP+FN)

Out[288]: 0.7961842717543043

In [289]: # Calculating Specificity
TN/(TN+FP)

Out[289]: 0.7854671280276817
```

16. Making predictions on test set

```
In [335]: # Making predictions
y_pred_final['final_predicted'] = y_pred_final.Conversion_Prob.map(lambda x: 1 if x > 0.46 else 0)

In [336]: y_pred_final.head()

Out[336]:
```


	Converted	Conversion_Prob	final_predicted
0	1	0.996712	1
1	0	0.137945	0
2	0	0.717442	1
3	1	0.311448	0
4	1	0.731041	1

17. Checking for features


```
In [343]: ▶ pd.options.display.float_format = '{:.2f}'.format
new_params = res.params[1:]
new_params
```

```
Out[343]: TotalVisits 10.23
Total Time Spent on Website 4.40
Lead Origin_Lead Add Form 4.24
Lead Source_Olark Chat 1.63
Lead Source_Welingak Website 2.34
Do Not Email_Yes -1.52
Last Activity_Had a Phone Conversation 1.17
Last Activity_SMS Sent 1.18
What is your current occupation_Housewife 22.61
What is your current occupation_Student -1.13
What is your current occupation_Unemployed -1.30
What is your current occupation_Working Professional 1.25
Last Notable Activity_Had a Phone Conversation 23.01
Last Notable Activity_Unreachable 2.77
Specialization_Select -0.34
dtype: float64
```

18. selection of features



```
In [353]: ▶ pd.DataFrame(feature_importance).reset_index().sort_values(by=0,ascending=False).head(3)
```

Out[353]:

	index	0
12	Last Notable Activity_Had a Phone Conversation	100.00
8	What is your current occupation_Housewife	98.26
0	TotalVisits	44.48